

S3 Task

1) Create an S3 bucket and upload some objects to S3

Go to **AWS Console** → **S3**

Click **Create bucket**

Enter a bucket name: `mybucket-techie`


Choose a region (example: EU-North-1)

Keep all default settings

Click **Create bucket**


—> s3 —> create bucket

—> Created a bucket name mybucket-techie



—> went inside the bucket and —>upload —> add file

—> after adding the file —>click upload



Upload succeeded

For more information, see the [Files and folders](#) table.



Upload: status

Close

After you navigate away from this page, the following information is no longer available.

Summary

Destination

s3://my-vpc-bucket-flowlog

Succeeded

1 file, 1.6 MB (100.00%)

Failed

0 files, 0 B (0%)

Files and folders

Configuration

Files and folders (1 total, 1.6 MB)

Find by name

< 1 >

Name	Folder	Type	Size	Status	Error
peakpx-2.jpeg	-	image/jpeg	1.6 MB	Succeeded	-

—> we have successfully added the object inside our mybucket-techie

2) Deploy a static website in the S3 bucket.

—> go inside the bucket which we created

—> now in side the properties scroll down until you see static website

The screenshot shows the AWS Management Console for the 'mybucket-techie' S3 bucket. The left sidebar contains navigation links for Amazon S3, Buckets, and various management and security settings. The main content area displays the bucket's properties, including Transfer acceleration (disabled), Object Lock (disabled), Requester pays (disabled), and Static website hosting (disabled). A blue banner at the bottom of the Static website hosting section recommends using AWS Amplify Hosting and provides a button to 'Create Amplify app'.

hosting

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

Edit

We recommend using AWS Amplify Hosting for static website hosting

Deploy a fast, secure, and reliable website quickly with AWS Amplify Hosting. Learn more about [Amplify Hosting](#) or [View your existing Amplify apps](#)

Create Amplify app

S3 static website hosting

Enabled

Hosting type

Bucket hosting

Bucket website endpoint

When you configure your bucket as a static website, the website is available at the AWS Region-specific website endpoint of the bucket. [Learn more](#)

<http://mybucket-technie-s3-website.eu-north-1.amazonaws.com>

>enable it

Inside the bucket → click **Properties**

Scroll down to **Static website hosting**


Click **Edit**

Enable it

Choose **Host a static website**

Set:

- Index document → `index.html`
- Error document → `error.html`



Save

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

Static website hosting

☐ Disable

☒ Enable

Hosting type

☒ Host a static website

Use the bucket endpoint as the web address. [Learn more](#)

☐ Redirect requests for an object

Redirect requests to another bucket or domain. [Learn more](#)

For your customers to access content at the website endpoint, you must make all your content publicly readable. To do so, you can edit the S3 Block Public Access settings for the [Public Access](#)

Index document

Specify the home or default page of the website.

index.html

Error document - optional

This is returned when an error occurs.

error.html

Redirection rules - optional


Redirection rules, written in JSON, automatically redirect webpage requests for specific content. [Learn more](#)

1		
---	--	--

—> now create this files and object this files in bucket

```
connection to cat 51 151 15 151:connect to amazonaws.com
waseemakram@waseem downloads % vi index.html
waseemakram@waseem downloads % cat index.html
hello world
waseemakram@waseem downloads % vi error.html
waseemakram@waseem downloads % cat error.html
site under maintenance...
waseemakram@waseem downloads %
```


- given bucket policy for static website




—>“Block Public Access”

If you haven't done this:


1. Go to **Permissions**
2. At the top → **Block public access (bucket settings)**
3. Click **Edit**
4. Turn **OFF**
 - “Block all public access”
5. Save

Without this, your website will not open.




—> after uploading index.html and error.html


—> scroll to the static website url and copy the url and paste in browser



- > we can see index.html page is displaying
- > so if we remove the index.html file from objects of bucket
- > it will show the error.html file on the same static website site



- > we can see it is showing error.html file



3) Enable cross-region replication on S3 buckets.


CRR = copying objects automatically from **one bucket** to another **in a different region**.

—> we have a bucket name “mybucket-techie” in eu-north-1 (stockholm)


—> another bucket in virginia named "myvirginia-bucket1" US East (N. Virginia) us-east-1

—> Enable Versioning on Source Bucket

1. Open your **mybucket-techie** (your website bucket) from eu-north-1
2. Go to **Properties**
3. Find **Versioning**
4. Click **Edit**
5. Turn ON **Versioning**
6. Save




—>this are our two buckets from different regions



—>Enable Versioning on Destination Bucket

Open your destination bucket → Properties → Versioning → Enable.

—>Below is my destination bucket from N.virginia us-east-1 “myvirginia-bucket1”



—> Set Up CRR Rule

Go to your **source bucket** →”mybucket-techie”

Top tabs → click **Management**

1. Click **Create replication rule**
2. Name it anything (e.g., `replicate-to-us-east`)
3. Choose:
 - **Source:** This bucket “mybucket-techie”
 - —> click **Apply to all objects in the bucket**
4. Scroll → **Destination bucket**
5. Choose the bucket you created in another region (myvirgina-bucket1)
6. IAM Role: click **Create new role**
7. Leave everything default
8. Save rule

The screenshot shows the AWS Management Console interface for creating a new replication rule. The breadcrumb navigation at the top reads: Amazon S3 > Buckets > mybucket-techie > Replication rules > Create replication rule. The main heading is 'Create replication rule' with an 'Info' link. The configuration is divided into three sections: 1. 'Replication rule configuration' containing 'Replication rule name' (set to 'replicate-to-virginia'), 'Status' (radio buttons for 'Enabled' and 'Disabled', with 'Enabled' selected), and 'Priority' (set to '0'). 2. 'Source bucket' containing 'Source bucket name' (set to 'mybucket-techie'), 'Source Region' (set to 'Europe (Stockholm) eu-north-1'), and 'Choose a rule scope' (radio buttons for 'Limit the scope of this rule using one or more filters' and 'Apply to all objects in the bucket', with the latter selected). 3. 'Destination' section, which is partially visible at the bottom of the screen.

Destination

Destination
You can replicate objects across buckets in different AWS Regions (Cross-Region Replication) or you can replicate objects across buckets in the same AWS Region (Same-Region Replication). You can also specify a different bucket for each rule in the configuration. [Learn more](#)

☒ Choose a bucket in this account
☐ Specify a bucket in another account

Bucket name
Choose the bucket that will receive replicated objects.

myvirginia-bucket1 [Browse S3](#)

Destination Region
US East (N. Virginia) us-east-1

IAM role
Permission to access the specified resources

☒ Create new role
☐ Choose from existing IAM roles
☐ Enter IAM role ARN

—> we have successfully created the replication rule for both buckets from diff regions

Amazon S3 > Buckets > mybucket-techie > Replication rules

Replication configuration successfully updated
If changes to the configuration aren't displayed, choose the refresh button. Changes apply only to new objects. To replicate existing objects with this configuration, choose **Create replication job**.

Replication rules [Info](#)
Replication enables automatic and asynchronous copying of objects across buckets in the same or different AWS Regions. A replication configuration is a set of rules that define what options should be applied to a group of objects during replication.

Replication configuration settings
Configuration settings affect all replication rules in the bucket.

Source bucket
mybucket-techie

Source Region
Europe (Stockholm) eu-north-1

IAM role
[sScr_role_for_mybucket-techie](#)

Replication rules (1) [View details](#) [Edit rule](#) [Delete](#) [Actions](#) [Create replication rule](#)

Use replication rules to define options you want Amazon S3 to apply during replication such as server-side encryption, replica ownership, transitioning replicas to another storage class, and more. [Learn more](#)


Replication rule name	Status	Destination bucket	Destination Region	Priority	Scope	Storage class	Replica owner	Replication Time Control	KMS-encrypted objects (SSE-KMS or DSSE-KMS)	Replica modification sync
replicate-to-virginia	Enabled	s3://myvirginia-bucket1	US East (N. Virginia) us-east-1	0	Entire bucket	Same as source	Same as source	Disabled	Do not replicate	Disabled

—> now Now test:


1. Upload any new object to **source bucket**
2. Go check the **destination bucket**
3. You will see the object replicated automatically

Note: Replication works **only for new uploads after the rule is created**, not old files.

—> I have uploaded a new file named my file in source bucket and now checked the other bucket



—> so in my other bucket from diff region (destination bucket) we can see we are having myfile named file in side our destination bucket also



4) Configure a bucket policy so only the Admin user can see the objects of the S3 bucket.

—>What This Policy Will Do


- **Only the Admin IAM user** can access the bucket
- **Everyone else (public, other IAM users, roles)** is denied
- Admin user = IAM user ARN you specify

—>since I have **devops** IAM user already has **S3 Full Access**, we will make this user the **ONLY** user who can access your S3 bucket.

The screenshot shows the AWS IAM console interface. The left sidebar contains navigation links for Identity and Access Management (IAM), Access management, and Access reports. The main content area displays the details for the 'devops' user. The 'Summary' section shows the user's ARN, console access status, and last sign-in. The 'Permissions' tab is active, showing a list of permissions policies attached to the user. The table lists five policies: AccessFriendBucket, AmazonEC2ReadOnlyAccess, AmazonS3FullAccess, EC2-Only-Stockholm, and IAMUserChangePassword. The 'Permissions boundary' section is currently empty.

Policy name	Type	Attached via
AccessFriendBucket	Customer inline	Inline
AmazonEC2ReadOnlyAccess	AWS managed	Group EC2-ReadOnly-Group
AmazonS3FullAccess	AWS managed	Directly
EC2-Only-Stockholm	Customer managed	Directly
IAMUserChangePassword	AWS managed	Directly


—> copy the ARN of the the user




—>Final Bucket Policy — ONLY devops user can access the bucket named “myvirginia-bucket1”

—>S3 → myvirginia-bucket1 → Permissions → Bucket Policy → Edit

—> the below policy will allow only devops name user to have access for bucket name “ myvirginia-bucket1”




—> in the corner right side of the pic we can see we are in side the iam user named devops which we have given the access to “myvirginia-bucket1”



—> the devops user is able to access the bucket

—> if we see I have login to another iam user named myuser iam unable to have access to the bucket “myvirginia-bucket1”
—> as it is giving only access to the user devops



5) Set up lifecycle policies to automatically transition or delete objects based on specific criteria.

—> by this task we can delete objects in a particular bucket to be deleted based on given days

What a Lifecycle Rule Does

You can tell S3 to:

- Move objects to cheaper storage (Standard-IA, Glacier)
- Delete old objects automatically
- Delete previous versions (if versioning on)
- Expire incomplete uploads

Example rule we will create:

✓ After 30 days → Move to Standard-IA

✓ After 90 days → Move to Glacier

✓ After 365 days → Delete the object

This is the most common lifecycle rule used in projects.

Steps to do

buckets —> mybucket-techie —> lifecycle confi —> create lifecycle rule

The screenshot shows the AWS Management Console interface for creating a lifecycle rule. The breadcrumb trail at the top reads: Amazon S3 > Buckets > mybucket-techie > Lifecycle configuration > Create lifecycle rule. The main heading is 'Create lifecycle rule' with an 'Info' link. Below this is the 'Lifecycle rule configuration' section. It includes a 'Lifecycle rule name' field with a placeholder 'Enter rule name' and a note 'Up to 255 characters'. Under 'Choose a rule scope', the radio button 'Limit the scope of this rule using one or more filters' is selected. The 'Filter type' section explains that filters can be based on prefix, object tags, object size, or a combination. It includes a 'Prefix' field with a placeholder 'Enter prefix' and a note: 'Don't include the bucket name in the prefix. Using certain characters in key names can cause problems with some applications and protocols. [Learn more](#)'. The 'Object tags' section has an 'Add tag' button. The 'Object size' section has two checkboxes: 'Specify minimum object size' and 'Specify maximum object size', both of which are currently unchecked. A 'Learn more' link is also present at the bottom of the configuration section.

—> given name as “delete-old-objects”

Amazon S3 > Buckets > mybucket-techie > Lifecycle configuration > Create lifecycle rule

Lifecycle rule name

delete-old-objects

Up to 255 characters

Choose a rule scope

☐ Limit the scope of this rule using one or more filters

☒ Apply to all objects in the bucket

⚠

Apply to all objects in the bucket

If you want the rule to apply to specific objects, you must use a filter to identify those objects. Choose "Limit the scope of this rule using one or more filters". [Learn more](#)

☒ I acknowledge that this rule will apply to all objects in the bucket.

Lifecycle rule actions

Choose the actions you want this rule to perform.

☐ Transition current versions of objects between storage classes

This action will move current versions.

☐ Transition noncurrent versions of objects between storage classes

This action will move noncurrent versions.

☐ Expire current versions of objects

☐ Permanently delete noncurrent versions of objects

☐ Delete expired object delete markers or incomplete multipart uploads

These actions are not supported when filtering by object tags or object size.

Review transition and expiration actions

Current version actions

Day 0

No actions defined.

Noncurrent versions actions

Day 0

No actions defined.

Choose Rule Scope

Select:

✓ **Apply to all objects in the bucket**

(Do NOT select prefix unless needed.)

—>Click this checkbox:

✓ Expire current versions of objects

—> give days like delete objects after 30 days

—> we have given 30

Lifecycle rule actions
Choose the actions you want this rule to perform.

- ☐ Transition current versions of objects between storage classes
This action will move current versions.
- ☐ Transition noncurrent versions of objects between storage classes
This action will move noncurrent versions.
- ☒ Expire current versions of objects
- ☐ Permanently delete noncurrent versions of objects
- ☐ Delete expired object delete markers or incomplete multipart uploads
These actions are not supported when filtering by object tags or object size.

Expire current versions of objects
For version-enabled buckets, Amazon S3 adds a delete marker and the current version of an object is retained as a noncurrent version. For non-versioned buckets, Amazon S3 permanently deletes the object.

Days after object creation
30

Review transition and expiration actions

Current version actions

Day 0

- Objects uploaded

↓

Day 30

- Objects expire

Noncurrent versions actions

Day 0

No actions defined.

—> click create

Amazon S3 > Buckets > mybucket-techie > Lifecycle configuration

🔍 ⓘ

🟢 The rule "delete-old-objects" has been successfully added and the lifecycle configuration has been updated. It may take some time for the configuration to be updated. Refresh the lifecycle rules list if changes to the configuration aren't displayed. ✕

Lifecycle configuration
To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Default minimum object size for transitions
All storage classes 128K

Lifecycle rules (1) [View details](#) [Edit](#) [Delete](#) [Actions](#) [Create lifecycle rule](#)

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

🔍 Find lifecycle rules by name

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete actions	Incomplete multipart uploads
delete-old-objects	🟢 Enabled	Entire bucket	Expires	-	-	-

ObjectsMetadataPropertiesPermissionsMetricsManagementAccess Points

Lifecycle configuration

To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Default minimum object size for transitions
All storage classes 128K

Lifecycle rules (1)

View detailsEditDeleteActionsCreate lifecycle rule

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete mar...	Incomplete multipart upl...
☐ delete-old-objects	☒ Enabled	Entire bucket	Expires	-	-	-

[View lifecycle configuration](#)

Replication rules (1)

View detailsEdit ruleDeleteActionsCreate replication rule

Use replication rules to define options you want Amazon S3 to apply during replication such as server-side encryption, replica ownership, transitioning replicas to another storage class, and more. [Learn more](#)

Replication rule name	Status	Destination bucket	Destination Region	Priority	Scope	Storage class	Replica owner	Replication Time Control	KMS-encrypted objects (SSE-KMS or DSSE-KMS)	Replica modification sync
☐ replicate-to-virginia	☒ Enabled	s3://myvirginia-bucket1	US East (N. Virginia) us-east-1	0	Entire bucket	Same as source	Same as source	Disabled	Do not replicate	Disabled

[View replication configuration](#)

—> WHAT YOU JUST CREATED

You have created a lifecycle rule where:

- ✓ All objects in the bucket
- ✓ Will be deleted automatically
- ✓ After 30 days

This is a valid and complete lifecycle policy.

6) Push some objects to S3 using the AWS CLI.

—> aws configure
Configure your aws access

```
waseemakram@waseem downloads % aws configure
AWS Access Key ID [*****EQMJ]:
AWS Secret Access Key [*****H4pd]:
Default region name [eu-north-1]:
Default output format [json]:
```

—> aws s3 ls (check your s3 buckets from aws)

```
waseemakram@waseem downloads % aws s3 ls
2025-12-03 17:23:38 mybucket-techie
2025-12-03 16:58:32 myvirginia-bucket1
```

—> aws s3 cp ~/Downloads/devops s3://mybucket-techie/ --recursive
aws s3 cp ~/Downloads/devops s3://mybucket-techie/ --recursive


```
waseemakram@waseem downloads % aws s3 cp ~/Downloads/devops s3://mybucket-techie/ --recursive
```

```
upload: devops/bash scripting/Bash_scripting_notes to s3://mybucket-techie/bash scripting/Bash_scripting_notes
upload: devops/bash scripting/Bash_scripting_notes (1) to s3://mybucket-techie/bash scripting/Bash_scripting_notes (1)
upload: devops/bash scripting/Exit+status.txt to s3://mybucket-techie/bash scripting/Exit+status.txt
upload: devops/.DS_Store to s3://mybucket-techie/.DS_Store
upload: devops/bash scripting/Basic+Shell+Scripts.doc to s3://mybucket-techie/bash scripting/Basic+Shell+Scripts.doc
upload: devops/VPC (1).docx to s3://mybucket-techie/VPC (1).docx
upload: devops/bash scripting/Unix+Shell+Scripting.docx to s3://mybucket-techie/bash scripting/Unix+Shell+Scripting.docx
upload: devops/bash scripting/case+Scripts.doc to s3://mybucket-techie/bash scripting/case+Scripts.doc
upload: devops/bash scripting/for+loop+Scripts.doc to s3://mybucket-techie/bash scripting/for+loop+Scripts.doc
upload: devops/bash scripting/do-while+Scripts.doc to s3://mybucket-techie/bash scripting/do-while+Scripts.doc
upload: devops/bash scripting/if-then+Scripts.doc to s3://mybucket-techie/bash scripting/if-then+Scripts.doc
upload: devops/devops linux/Linux-commands to s3://mybucket-techie/devops linux/Linux-commands
upload: devops/devops linux/how to create user and add to user in sudoer fiel in linux to s3://mybucket-techie/devops linux/how to create user and
upload: devops/devops linux/Linux-Cheat-Sheet.png to s3://mybucket-techie/devops linux/Linux-Cheat-Sheet.png
upload: devops/devops linux/Linux_file_and_directory_management_1754976493.pdf to s3://mybucket-techie/devops linux/Linux_file_and_directory_manag
upload: devops/devops linux/linux-sheet1_0.jpg to s3://mybucket-techie/devops linux/linux-sheet1_0.jpg
upload: devops/devops linux/CRONTAB-FOR-JOB-SCHEDULING.pdf to s3://mybucket-techie/devops linux/CRONTAB-FOR-JOB-SCHEDULING.pdf
upload: devops/devops linux/Linux-Cheat-Sheet-Sponsored-By-Loggly.pdf to s3://mybucket-techie/devops linux/Linux-Cheat-Sheet-Sponsored-By-Loggly.p
upload: devops/devops linux/Linux-sheet2_0.jpg to s3://mybucket-techie/devops linux/Linux-sheet2_0.jpg
upload: devops/ec2 documents/EC2 Instance Types _ A Complete AWS EC2 Instance Information.pdf to s3://mybucket-techie/ec2 documents/EC2 Instance T
upload: devops/git hub /.DS_Store to s3://mybucket-techie/git hub /.DS_Store
upload: devops/git hub /Adding a GitHub Webhook in Your Jenkins Pipeline - DZone DevOps (1).pdf to s3://mybucket-techie/git hub /Adding a GitHub W
upload: devops/git hub /Git-Commands to s3://mybucket-techie/git hub /Git-Commands
upload: devops/git hub /Git-Commands (1) to s3://mybucket-techie/git hub /Git-Commands (1)
upload: devops/git hub /Git_Commands_1744427584.pdf to s3://mybucket-techie/git hub /Git_Commands_1744427584.pdf
upload: devops/git hub /Github task 1 (2).docx to s3://mybucket-techie/git hub /Github task 1 (2).docx
upload: devops/git hub /git important commands to learn to s3://mybucket-techie/git hub /git important commands to learn
upload: devops/git hub /How To Install Git Bash On Windows - Stanley Ulili (1).pdf to s3://mybucket-techie/git hub /How To Install Git Bash On Win
upload: devops/app web server& ha proxy/app-server web-serwer-task.docx to s3://mybucket-techie/app web server& ha proxy/app-server web-serwer-ta
upload: devops/ec2 documents/How to Create EC2 Instance in AWS_ Step by Step Tutorial.pdf to s3://mybucket-techie/ec2 documents/How to Create EC2
upload: devops/software-development-methodologies.docx to s3://mybucket-techie/software-development-methodologies.docx
upload: devops/app web server& ha proxy/AWS Task.docx to s3://mybucket-techie/app web server& ha proxy/AWS Task.docx
upload: devops/devops linux/software-development-methodologies (1).docx to s3://mybucket-techie/devops linux/software-development-methodologies (1)
upload: devops/app web server& ha proxy/Configure HA Proxy server (1).docx to s3://mybucket-techie/app web server& ha proxy/Configure HA Proxy ser
waseemakram@waseem downloads %
```



```
aws s3 cp ~/Downloads/devops s3://mybucket-techie/ --recursive
```

—> this will copy the devops name folder and will store in mybucket-techie named bucket



—> we are able to see the devops folder files in the “mybucket-techie” Bucket

—> aws s3 ls s3://mybucket-techie/

To check the objects from the bucket

```
upload: devops/app web server& ha proxy/configure ha proxy server (1).docx to s3://mybucket-techie/
waseemakram@waseem downloads % aws s3 ls s3://mybucket-techie/

      PRE app web server& ha proxy/
      PRE bash scripting/
      PRE devops linux/
      PRE ec2 documents/
      PRE git hub /
2025-12-03 18:59:52      14340 .DS_Store
2025-12-03 15:34:41    101893 Screenshot 2025-12-03 at 3.02.50 PM.png
2025-12-03 16:59:23    585996 Screenshot 2025-12-03 at 3.58.40 PM.png
2025-12-03 18:59:52     28010 VPC (1).docx
2025-12-03 14:52:22        27 error.html
2025-12-03 15:36:54         7 file9
2025-12-03 15:10:56        12 index.html
2025-12-03 18:59:55     91700 linux-and-itsm-and-patchmanagement-and-commands (1).docx
2025-12-03 15:32:38        15 myfile
2025-12-03 14:41:07    1651414 peakpx-2.jpeg
2025-12-03 18:59:55    26147906 software-development-methodologies.docx
waseemakram@waseem downloads %
```

7) Write a Bash script to create an S3 bucket.

```
waseemakram@waseem downloads % vi create_bucket.sh

[waseemakram@waseem downloads % chmod +x create_bucket.sh

[waseemakram@waseem downloads % ./create_bucket.sh

Creating S3 bucket: mybucket-32369 in region: eu-north-1
{
  "Location": "http://mybucket-32369.s3.amazonaws.com/"
}
Bucket created successfully: mybucket-32369
waseemakram@waseem downloads %
```

—> created a file name create_bucket.sh

—> and added the below script for creating the new bucket

```
#!/bin/bash

# Generate a unique bucket name using random number
BUCKET_NAME="mybucket-$RANDOM"
REGION="eu-north-1"


echo "Creating S3 bucket: $BUCKET_NAME in region: $REGION"

aws s3api create-bucket \
  --bucket $BUCKET_NAME \
  --region $REGION \
  --create-bucket-configuration LocationConstraint=$REGION

echo "Bucket created successfully: $BUCKET_NAME"

~
~
~
~
~
~
~
~
~
~
```

And after running the script we are able to see the bucket is created “mybucket-32369”



The screenshot shows the AWS Management Console interface for 'General purpose buckets'. At the top, there are tabs for 'General purpose buckets' (selected) and 'Directory buckets'. Below the tabs, there's a section titled 'General purpose buckets (3)' with an 'Info' link. To the right of this title are buttons for 'Copy ARN', 'Empty', 'Delete', and a prominent orange 'Create bucket' button. Below the title, a subtitle states 'Buckets are containers for data stored in S3.' A search bar with the placeholder 'Find buckets by name' is present. Below the search bar is a table with three columns: 'Name', 'AWS Region', and 'Creation date'. The table lists three buckets: 'mybucket-32369' in the 'Europe (Stockholm) eu-north-1' region, 'mybucket-techie' in the 'Europe (Stockholm) eu-north-1' region, and 'myvirginia-bucket1' in the 'US East (N. Virginia) us-east-1' region. Each bucket name is a blue hyperlink. The creation dates are listed as 'December 3, 2025, 19:08:16 (UTC+05:30)', 'December 3, 2025, 14:35:14 (UTC+05:30)', and 'December 3, 2025, 15:22:08 (UTC+05:30)' respectively.

Name	AWS Region	Creation date
mybucket-32369	Europe (Stockholm) eu-north-1	December 3, 2025, 19:08:16 (UTC+05:30)
mybucket-techie	Europe (Stockholm) eu-north-1	December 3, 2025, 14:35:14 (UTC+05:30)
myvirginia-bucket1	US East (N. Virginia) us-east-1	December 3, 2025, 15:22:08 (UTC+05:30)

—>aws s3 ls

To check the s3 bucket list

```
waseemakram@waseem downloads % aws s3 ls
2025-12-03 19:08:16 mybucket-32369
2025-12-03 17:23:38 mybucket-techie
2025-12-03 16:58:32 myvirginia-bucket1
waseemakram@waseem downloads %
```

8) Upload a 1 GB file to S3 using the CLI.

—> I have a 1gb video were I have uploaded it in side bucket name “mybucket-techie”

—>aws s3 cp my1gbvideo.mov s3://mybucket-techie/

This will copy the video and will push it to the “mybucket-techie”

```
The user-provided path my1gbvideo does not exist.
waseemakram@waseem Desktop % aws s3 cp my1gbvideo.mov s3://mybucket-techie/

Completed 122.0 MiB/1.6 GiB (6.5 MiB/s) with 1 file(s) remaining
```

—> successfully uploaded the video to bucket

```
The user-provided path my1gbvideo does not exist.
waseemakram@waseem Desktop % aws s3 cp my1gbvideo.mov s3://mybucket-techie/

upload: ./my1gbvideo.mov to s3://mybucket-techie/my1gbvideo.mov
waseemakram@waseem Desktop %
```

—> now go and check that bucket we are able to see the 1.6gb object is available in side the bucket

ttings	☐	linux-and-itsm-and-patchmanagement-and-commands (1).docx	docx	December 3, 2025, 18:59:55 (UTC+05:30)	89.6 KB	Standard
	☐	my1gbvideo.mov	mov	December 3, 2025, 19:23:40 (UTC+05:30)	1.6 GB	Standard
	☐	myfile	-	December 3, 2025, 15:32:38 (UTC+05:30)	15.0 B	Standard
	☐	peakpx-2.jpeg	jpeg	December 3, 2025, 14:41:07 (UTC+05:30)	1.6 MB	Standard

